

Exploiting Spectral Augmentation for Code-Switched Spoken Language Identification

Pradeep Rangan*, Sundeeep Teki*, and Hemant Misra

Applied Research Department, Swiggy (Bundl Technologies India Limited)

[pradeep.rl, sundeeep.teki, hemant.misra] @swiggy.in

Abstract

Spoken language Identification (LID) systems are needed to identify the language(s) present in a given audio sample, and typically could be the first step in many speech processing related tasks such as automatic speech recognition (ASR). Automatic identification of the languages present in a speech signal is not only scientifically interesting, but also of practical importance in a multilingual country such as India. In many of the Indian cities, when people interact with each other, as many as three languages may get mixed. These may include the official language of that province, Hindi and English (at times the languages of the neighboring provinces may also get mixed during these interactions). This makes the spoken LID task extremely challenging in Indian context. While quite a few LID systems in the context of Indian languages have been implemented, most such systems have used small scale speech data collected internally within an organization. In the current work, we perform spoken LID on three Indian languages (Gujarati, Telugu, and Tamil) code-mixed with English. This task was organized by the Microsoft research team as a spoken LID challenge. In our work, we modify the usual spectral augmentation approach and propose a language mask that discriminates the language ID pairs, which leads to a noise robust spoken LID system. The proposed method gives a relative improvement of approximately 3-5% in the LID accuracy over a baseline system proposed by Microsoft on the three language pairs for two shared tasks suggested in the challenge.

Index Terms: spoken language identification, code-mixing, spectral augmentation, noise-robustness, speech recognition

1. Introduction

In recent years, advances in artificial intelligence (AI) have significantly expanded the degree to which individuals can interact with technology utilizing just their voice [1]. However, these systems often require explicit information about the language of the users. The ability to dynamically process a single input speech stream consisting of different languages would expand the usefulness of existing voice-based systems and open up a wide array of additional functionalities. The spoken language identification (LID) research addresses this issue by exploring how to extract information from the audio signal and use it to predict the spoken language. LID is used in several applications such as multilingual translation systems or emergency call routing, where the response time of a fluent native operator might be critical [2].

Over the years, researchers have utilized many prosodic and acoustic features to construct machine learning models for LID systems [3]. Several **prosodic** and acoustic features are based on phonemes, which become the underlying features that drive the

performance of the statistical models [4]. **If two languages have many overlapping phonemes, then identifying them becomes a challenging task for a classifier. Subsequent LID systems relied on acoustic modelling [5].** In particular, guided by the advances on speaker verification, **the use of i-vector feature extractors as a front-end followed by diverse classification mechanisms became popular as acoustic LID systems [6] [7].** The extensive feature engineering with i-vectors results in very complex systems, with an increasing number of computational steps in their pipeline [8].

Approaches solely based on applying neural networks on input features like mel-frequency cepstral coefficients (MFCC) show that they reach state-of-the-art results, while being less complex [9]. Current research on language identification systems using deep neural networks (DNN) mainly focuses on using different forms of long short term memory (LSTMs), working on **input sequences of transformed audio data.** The resulting sequence features are fused together and used to classify the language of the input samples. **In [10] [11], a DNN based architecture was proposed for extracting spatial features from the log-mel spectrograms of raw audio using convolutional neural networks (CNNs) and then using recurrent neural networks (RNNs) for capturing temporal features to identify the language.**

Automatic spoken LID systems are particularly relevant for multilingual countries such as India. **An LID system recognising 27 Indian languages was implemented using Gaussian mixture model (GMM) and MFCC features [12].** Another LID system for Bengali, Hindi, Telugu, Urdu, Assamese, Punjabi and Manipuri languages was implemented that used feed forward neural networks trained with two hours of speech data of each of the seven languages [13]. Most LID systems for Indian languages were trained and tested with either speech of professional news readers [2] or with small scale speech data collected by researchers [14] [15], that may not be available for public use. Till recently, the lack of publicly available Indian language speech databases was a limitation for a scientific study of various acoustic analysis and modeling techniques. **Recently, a phonetically balanced speech corpus of Hindi-English code-mixed speech was developed by IIT Guwahati [16], and is explicitly designed for automatic speech recognition (ASR) paradigm, but not for the LID task.**

Microsoft organized a shared task on Code-switched spoken language identification (LID) in three language pairs – Gujarati-English, Telugu-English and Tamil-English [17]. The shared task consists of two subtasks: (1) utterance-level identification of monolingual vs. code-switched utterances, and, (2) frame-level identification of languages in a code-switched utterance.

Spectral augmentation is a popular technique to improve the noise robustness of an ASR system [18] [19]. It involves three main steps namely time warping, frequency masking, and tem-

*These authors contributed equally to this work

temporal masking. In this paper, we have explored the spectral augmentation [20] for detecting the LID information. The spectral augmentation approach randomly chooses the position of the different masks. However, in our work we select the temporal masks based on the number and position of the English segments in a code-switched utterance. The clean and the masked spectrograms are provided as an input to the CNN-BiLSTM encoder network trained with CTC loss function. The variation of frequency and temporal masks in the recent spectral augmentation [20] are found to provide discriminative power to identify the language information in a code-mixed speech data.

The organization of the remaining sections of the paper are as follows: In Section 2, the importance of spectral augmentation for LID is discussed. The proposed framework for spoken LID is explained in Section 3. Section 4 provides the experimental set-up of the proposed spectral augmentation based spoken LID. We analyze the impact of the proposed approach on spoken LID performance in Section 5. Section 6 describes the conclusions and future directions.

2. Spectral Augmentation

In this section, the related work on spectral augmentation for different speech processing applications is discussed. Data augmentation is a popular method for improving robustness and training of neural networks [18]. It's been used successfully in several domains starting from image classification to molecular modelling [21]. The fundamental principle is to increase the amount of training samples by creating multiple variants of the dataset. The data augmentation method is typically applied to create an additional training data for ASR to improve the performance. For instance, in [22], the data was augmented for low resource speech recognition tasks. Vocal Tract Length Normalization has been explored for data augmentation in [23]. Speech perturbation has been applied on raw audio for LVCSR tasks in [19]. The use of an acoustic room simulator has been adopted in [22]. Data augmentation is used to spot the important keywords from the speech utterance [24]. Perceptually, human listeners show remarkable tolerance to a variety of spectrotemporal manipulations of the input sound signal during segregation of speech-in-noise recognition tasks [25] [26] [27]. Inspired by such studies in human auditory perception and cognitive neuroscience, and the recent success of augmentation techniques in the speech and vision domains, SpecAugment, an augmentation method that operates on the log mel spectrogram of the input audio (rather than the raw audio) was proposed in [20]. This method is simple and computationally inexpensive, as it directly acts on the log mel spectrogram as if it was an image.

Spectral Augmentation includes the following steps:

1. **Time warping:** Given a log mel spectrogram with τ time steps, it is viewed as an image where horizontal and vertical axes represents time (s) and the frequency (Hz) respectively. A random point along the horizontal line passing through the center of the image within the time steps ($W, \tau - W$) is to be warped either to the left or right by a distance w chosen from a uniform distribution from 0 to the time warp parameter W along that line.
2. **Frequency Masking:** It is applied so that f consecutive mel frequency channels ($f_0, f_0 + f$) are masked, where f is first chosen from a uniform distribution from 0 to the frequency mask parameter F , and f_0 is chosen from $[0, v - f]$. v is the number of mel frequency channels.
3. **Temporal Masking:** It is applied so that t consecutive

time steps $[t_0, t_0 + t)$ are masked, where t is first chosen from a uniform distribution from 0 to the time mask parameter T , and t_0 is chosen from $[0, \tau - t)$.

Figure 1 shows examples of the individual augmentations applied to a single speech input (PartBGujarati/Dev/Audio/000060438.wav). The log mel spectrograms are normalized to have zero mean value, and thus setting the masked value to zero is equivalent to setting it to the mean value.




Figure 1: Augmentations applied to the base input. The figures (from top) depict the log mel spectrogram of the base input with no augmentation, (time warp, frequency and time masking applied)

The spectral augmentation approach is shown to be successful on Noise robust ASR [18]. For instance, the selection of different masks in the spectral augmentation can discriminate between the clean speech and the noisy speech. In code-mixed speech spoken in India, it is generally observed that the native speaker tries to speak in her native language, and occasionally uses English words. The process of masking the English words can lead to building a mono-lingual corpus. In order to build an efficient spoken LID system, the DNN may first learn to discriminate between the code-mixed speech and the mono-lingual speech. Since the temporal masking positions can be altered with respect to different languages, we sought to mask out the English words in the code-switched speech.

3. Proposed Work

Figure 2 shows the CNN-LSTM system that uses CTC loss function at the output layer. The fundamental end-to-end CTC




Figure 2: End-to-End spoken LID using CNN-LSTM model

system receives the 2-D spectrogram as shown in the bottom of the figure. The CNN layer is effectively used in speech processing applications as they are capable of modeling temporal invariance for variable length utterances [28]. The BiLSTM

models are trained on the convolved features to capture long-term sequential context. Later, the dense layer or the fully connected (FC) layer is connected to the output layer. The output is produced through a softmax function computing a probability distribution over the target labels. The target labels used in Figure 2 are S, G and E which corresponds to Silence, Gujarati and English language ID, respectively. To accelerate the training procedure, Batch Normalization [1] is applied on hidden layers. In our work, we feed the clean spectrogram, and the augmented spectrogram to the end-to-end LID system. However, the spectral augmentation is obtained by proposing an appropriate temporal mask termed as ‘language mask’ in the paper, and explained in the next subsection.

3.1. Language Transcript based Temporal mask

The conventional temporal mask(s) used in the spectral augmentation is selected at a random position that follows a uniform distribution from 0 to the time mask parameter T . In our work, at the time of training, the temporal masks are obtained from the language ID transcripts. The steps followed to obtain the temporal mask from the language transcript are as follows:

1. Initially, the number of characters in the transcript is computed (l_1).
2. Find the language segments from the transcript. Ex: If transcript = ‘SSSGGEEGG’, then the language segments are ‘SSS’, ‘GGG’, ‘EE’, and ‘GG’, where ‘S’, ‘E’ and ‘G’ corresponds to the silence, English, and Gujarati speech segments, respectively.
3. Obtain the maximum number of frames based on the length of the characters in the transcript (N_f). Since the language label is created for every 200 ms for the shared task, every language ID character corresponds to 200 ms of speech content.
4. We calculate the number of temporal masks by finding the number of speech segments corresponding to the English speech. Since, the non-predominant language in the code-mixed speech data is English, we mask the English words from the spectrogram.
5. Obtain the start (t_0), and end time frame index (t_1) for every English speech segment.
6. The t_0 to t_1 consecutive time steps $[t_0, t_1]$ are masked.

Figure 3 shows a toy example to augment the spectrogram based on the language mask. The spectrogram is obtained for the speech signal ‘PartB/Gujarati/000010183.wav’ as shown in Figure 3 (a). The speech signal is sampled at 16 KHz with window size of 200 ms, and window stride of 100 ms. The language transcripts are mapped to the speech frames, and are shown in Figure 3 (b). For instance, there is an English phrase (‘flight operation’) in the speech signal from 3.21 s to 3.99 s (shown in black). The number of temporal masks are decided based on the number of English segments present in the language transcripts. In this case, the number of English segments are two, and are present from (a) 3.1 - 3.99 s, and (b) 4.69 - 4.79 s. The corresponding time segments are masked as shown in Figure 3 (c).

4. Experiments

First, we introduce our experimental environment and the metrics used for the spoken LID task. Then, we show our results on the shared tasks. Following this, we show the results of our experiments on the proposed work.

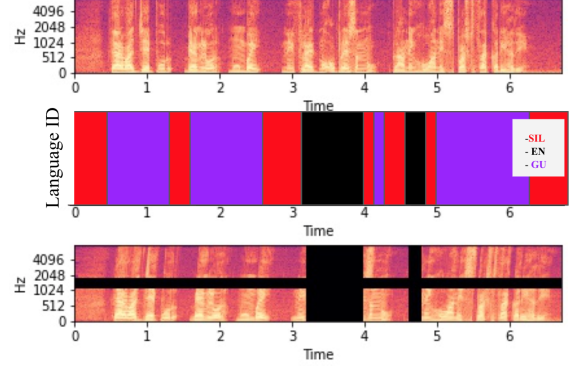


Figure 3: Illustrative example of the proposed LID based temporal masking in spectral augmentation: (Top) Raw spectrogram, (Middle) Language information at different time instance, and (Bottom) Proposed spectral augmentation

4.1. Data

The data set consists of three code-switched language pairs - Gujarati-English, Tamil-English and Telugu-English. The shared task for the workshop consists of two sub-tasks for spokenLID of code-switched audio. The two sub-tasks will consist of (1) Part A: Utterance-level LID, and (2) Part B: Frame-level LID.

Table 1: Dataset description for the code-mixed spoken LID

Data	Language	Number of Samples (Duration in hh:mm:ss)		
		Train	Dev	Test
Part-A	Gu-En	16780 (31:59:53)	2091 (3:59:55)	2156 (3:44:50)
	Ta-En	17943 (31:59:47)	2258 (4:00:05)	2258 (3:53:59)
	Te-En	16991 (31:59:50)	2135 (4:00:06)	2064 (3:42:15)
Part-B	Gu-En	8620 (15:59:56)	1080 (1:59:58)	1078 (2:00:02)
	Ta-En	8982 (15:59:57)	1135 (1:59:59)	1129 (1:59:57)
	Te-En	8226 (15:59:56)	1047 (2:00:03)	1033 (1:59:55)

4.2. Training

The Microsoft’s baseline system is made up of an end-to-end multi-layer model consisting of 5 layers of LSTM, each consisting of 1024 neurons. The model is based on deepspeech-2 [29]. The model is trained using the CTC loss function. The LID detector using CTC starts with two layers of 2D convolutions over both time and frequency domains with 32 channels, 41×11 , 21×11 filter dimensions, and 2×2 , 2×1 stride. Next, five Bi-LSTM layers with 1024 hidden units are followed by one fully connected linear layer with 5 softmax outputs blank, , Gujarati, English, Silence. The Bi-LSTM models have around 10.2 millions (M) parameters. The input sequence are values of spectrogram slices, 20 ms long, computed from Hamming windows with 10 ms frame shifts. The output (target) sequence was obtained directly from the letters of the word transcription. We used 100 epochs to train all the models used for further evaluation. The two variants of spectral augmentation are implemented in this paper: (1) spectral augmentation with random temporal, and frequency mask positions, and (2) spectral augmentation with the proposed language mask.

4.3. Decoding

Assuming an input sequence of length T , the output of the neural network will be $p(c|x_t)$ for $t = 1, \dots, T$. Let $p(c|x_t)$ is a distribution over possible characters in the alphabet ψ (which includes the blank symbol) given audio input x_t . In order to recover a character string from the output of the neural network, as a first approximation, we take the argmax at each time step. Let $S = (s_1, \dots, s_T)$ be the character sequence where

$$s_t = \arg \max_{c \in \psi} p(c|x_t) \quad (1)$$

The sequence S is mapped to a transcription by collapsing repeat characters and removing blanks. On the other hand, the beam search decoder uses the context information in generating the decoded sequence. There are two cases: either we extend the beam by a character c different from the last character, then there is no need for separating blanks in the paths, or the last character is repeated.

4.4. Evaluation

For task A, the predicted label (Monolingual or code-switched) file is submitted for a particular audio file in the blind test set. For task B, a frame-level (200ms) label for each frame in the audio is submitted. The LID system performance is evaluated using accuracy and Equal Error Rate (EER) as evaluation metrics.

$$Accuracy = \frac{N}{T} \quad (2)$$

where N and T are the total no. of correctly predicted data samples and the total no. of data points in the speech dataset respectively.

$$EER = \frac{FRR + FAR}{2} \quad (3)$$

where $FRR = TFR/T$, and $FAR = TFA/T$; FRR , FAR , TFR , TFA and T corresponds to false rejection rate, false acceptance rate, total no. of false rejects, total no. of false accepts and total number of datapoints, respectively.

5. Results and Discussion

The DeepSpeech-2 model is trained on the baseline configurations and on the variants of spectral augmentation. In this section, we report the accuracies and EERs on the Test set for Task-A and Task-B.

Table 2 shows the % spoken LID Accuracy of the baseline model and the proposed spectral augmentation method. It can be observed that the proposed system outperforms as compared to that of the baseline model on Task-A.

Table 2: *Spoken LID accuracy on Task-A for different datasets*

Data_Test	% Acc [Baseline]	% Acc [SpecAug]
Gu-En	71.9	73.01
Ta-En	71.2	79.02
Te-En	74.0	78.65

For Task-B, we compare the performance of the code-mixed language recognizer using spectral augmentation with random temporal mask location, and the proposed language mask. Table 3 shows the performance comparison of the proposed work, and the DeepSpeech-2 model (Microsoft’s Baseline) on the test set of Part-B with Greedy Search Decoding. It

Table 3: *Performance comparison of the proposed work, and the DeepSpeech-2 model (Microsoft’s Baseline) on the test set of Part-B by Greedy Search Decoder*

Language	% Acc [%EER] on the test set on Greedy search decoder		
	Baseline	SpecAug	SpecAug + Lang Mask
Gujarati	66.79 [9.71]	75.33 [7.78]	75.72 [7.53]
Tamil	72.17 [8.42]	74.80 [7.76]	75.02 [7.67]
Telugu	70.54 [8.65]	74.06 [7.88]	74.08 [7.87]

Table 4: *Performance comparison of the proposed work, and the DeepSpeech-2 model (Microsoft’s Baseline) on the test set of Part-B by Beam Search Decoder*

Language	% Acc [%EER] on the test set on Beam search decoder		
	Baseline	SpecAug	SpecAug + Lang Mask
Gujarati	72.53 [8.54]	76.33 [7.52]	76.64 [7.36]
Tamil	73.89 [8.06]	75.80 [7.54]	76.06 [7.44]
Telugu	75.20 [7.68]	75.90 [7.47]	75.84 [7.47]

can be observed that the proposed language mask on the spectrogram is able to discriminate the language ID’s within an utterance.

Table 4 shows the performance comparison of the proposed work, and the state-of-the-art methods on the test set of Part-B on Beam Search Decoder. The size of the beam width is varied from 5 to 20 in steps of 5. It is observed that the beam width size of 15 is optimal for all languages in the shared task.

5.1. Discussion

The research on spoken LID for code-switched speech is relatively new for Indian languages. Further, in India it is fairly common to mix words from English and Hindi along with the native language. Spectral augmentation has shown promising improvement on noise robust ASR, and on low resource languages. Instead of masking random positions in the spectrogram, the positions of the language transitions are masked. Due to this, the train speech corpus contains only the code-switched, and the mono-lingual utterances. As a result, the model learns to discriminate well between the code-switched utterances, and the mono-lingual utterances. Also, the positions of the proposed language mask captures the corresponding grammar in the utterance. Thereby the model learns about the different statistics of code-mixing. On the speech segments where there are small portions of non-dominant language is involved, the model captures these small language transitions.

6. Conclusion

Spoken LID is performed on the three Indian languages (Gujarati, Telugu, and Tamil) code-switched with English. We proposed a language mask from the speech transcripts, and incorporated it in the spectrogram. We observed that the proposed method is able to discriminate the languages in a code-mixed data. The proposed method gives a relative improvement in performance of approximately 3-5% in the LID accuracy over that of a baseline system proposed by Microsoft on three code-switched language pairs for two different shared tasks.

7. References

- [1] G. Hinton, L. Deng, D. Yu, G. E. Dahl, A.-r. Mohamed, N. Jaitly, A. Senior, V. Vanhoucke, P. Nguyen, T. N. Sainath *et al.*, “Deep neural networks for acoustic modeling in speech recognition: The

- shared views of four research groups,” *IEEE Signal processing magazine*, vol. 29, no. 6, pp. 82–97, 2012.
- [2] C. Kumar and H. Li, “Language identification for multilingual speech recognition systems,” in *9th Conference Speech and Computer*, 2004.
 - [3] Y. Obuchi and N. Sato, “Language identification using phonetic and prosodic hmms with feature normalization,” in *Proceedings (ICASSP’05). IEEE International Conference on Acoustics, Speech, and Signal Processing, 2005.*, vol. 1. IEEE, 2005, pp. I–569.
 - [4] N. E. Safitri, A. Zahra, and M. Adriani, “Spoken language identification with phonotactics methods on minangkabau, sundanese, and javanese languages,” in *SLTU*, 2016, pp. 182–187.
 - [5] R. Tong, B. Ma, D. Zhu, H. Li, and E. S. Chng, “Integrating acoustic, prosodic and phonotactic features for spoken language identification,” in *2006 IEEE International Conference on Acoustics Speech and Signal Processing Proceedings*, vol. 1. IEEE, 2006, pp. I–I.
 - [6] J. Gonzalez-Dominguez, I. Lopez-Moreno, J. Franco-Pedroso, D. Ramos, D. T. Toledano, and J. Gonzalez-Rodriguez, “Multi-level and session variability compensated language recognition: Atvs-uam systems at nist lre 2009,” *IEEE Journal of Selected Topics in Signal Processing*, vol. 4, no. 6, pp. 1084–1093, 2010.
 - [7] P. A. Torres-Carrasquillo, E. Singer, T. Gleason, A. McCree, D. A. Reynolds, F. Richardson, and D. Sturim, “The mitll nist lre 2009 language recognition system,” in *2010 IEEE International Conference on Acoustics, Speech and Signal Processing*. IEEE, 2010, pp. 4994–4997.
 - [8] D. Martinez, O. Plhot, L. Burget, O. Glembek, and P. Matějka, “Language recognition in ivectors space,” in *Twelfth annual conference of the international speech communication association*, 2011.
 - [9] G. Gelly, J.-L. Gauvain, L. Lamel, A. Laurent, V. B. Le, and A. Messaoudi, “Language recognition for dialects and closely related languages,” in *Odyssey*, vol. 2016, 2016, pp. 124–131.
 - [10] S. Shukla, G. Mittal *et al.*, “Spoken language identification using convnets,” in *European Conference on Ambient Intelligence*. Springer, 2019, pp. 252–265.
 - [11] C. Bartz, T. Herold, H. Yang, and C. Meinel, “Language identification using deep convolutional recurrent neural networks,” in *International Conference on Neural Information Processing*. Springer, 2017, pp. 880–889.
 - [12] S. G. Koolagudi, D. Rastogi, and K. S. Rao, “Identification of language using mel-frequency cepstral coefficients (mfcc),” *Procedia Engineering*, vol. 38, pp. 3391–3398, 2012.
 - [13] K. S. Rao, V. R. Reddy, and S. Maity, *Language identification using spectral and prosodic features*. Springer, 2015.
 - [14] C. Madhu, A. George, and L. Mary, “Automatic language identification for seven indian languages using higher level features,” in *2017 IEEE International Conference on Signal Processing, Informatics, Communication and Energy Systems (SPICES)*. IEEE, 2017, pp. 1–6.
 - [15] J. Chakraborty, S. Nath, S. Nirmala, and K. Samudravijaya, “Language identification of assamese, bengali and english speech,” in *SLTU*, 2018, pp. 177–181.
 - [16] A. Pandey, B. M. L. Srivastava, R. Kumar, B. T. Nellore, K. S. Teja, and S. V. Gangashetty, “Phonetically balanced code-mixed speech corpus for hindi-english automatic speech recognition,” in *Proceedings of the Eleventh International Conference on Language Resources and Evaluation (LREC 2018)*, 2018.
 - [17] Microsoft, *Spoken LID Challenge*, 2020. [Online]. Available: <https://www.microsoft.com/en-us/research/event/workshop-on-speech-technologies-for-code-switching-2020/>
 - [18] T. Ko, V. Peddinti, D. Povey, and S. Khudanpur, “Audio augmentation for speech recognition,” in *Sixteenth Annual Conference of the International Speech Communication Association*, 2015.
 - [19] L. Tóth, G. Kovács, and D. Van Compernelle, “A perceptually inspired data augmentation method for noise robust cnn acoustic models,” in *International Conference on Speech and Computer*. Springer, 2018, pp. 697–706.
 - [20] D. S. Park, W. Chan, Y. Zhang, C.-C. Chiu, B. Zoph, E. D. Cubuk, and Q. V. Le, “SpecAugment: A simple data augmentation method for automatic speech recognition,” *arXiv preprint arXiv:1904.08779*, 2019.
 - [21] E. J. Bjerrum, M. Ghlader, and T. Skov, “Data augmentation of spectral data for convolutional neural network (cnn) based deep chemometrics,” *arXiv preprint arXiv:1710.01927*, 2017.
 - [22] A. Ragni, K. Knill, S. P. Rath, and M. Gales, “Data augmentation for low resource languages,” 2014.
 - [23] N. Jaitly and G. E. Hinton, “Vocal tract length perturbation (vtlp) improves speech recognition,” in *Proc. ICML Workshop on Deep Learning for Audio, Speech and Language*, vol. 117, 2013.
 - [24] A. Raju, S. Panchapagesan, X. Liu, A. Mandal, and N. Strom, “Data augmentation for robust keyword spotting under playback interference,” *arXiv preprint arXiv:1808.00563*, 2018.
 - [25] S. Teki, M. Chait, S. Kumar, K. von Kriegstein, and T. D. Griffiths, “Brain bases for auditory stimulus-driven figure-ground segregation,” *Journal of Neuroscience*, vol. 31, no. 1, pp. 164–171, 2011.
 - [26] S. Teki, M. Chait, S. Kumar, S. Shamma, and T. D. Griffiths, “Segregation of complex acoustic scenes based on temporal coherence,” *Elife*, vol. 2, p. e00699, 2013.
 - [27] S. Teki, N. Barascud, S. Picard, C. Payne, T. D. Griffiths, and M. Chait, “Neural correlates of auditory figure-ground segregation based on temporal coherence,” *Cerebral cortex*, vol. 26, no. 9, pp. 3669–3680, 2016.
 - [28] Y. Miao, M. Gawayyed, X. Na, T. Ko, F. Metze, and A. Waibel, “An empirical exploration of ctc acoustic models,” in *2016 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2016, pp. 2623–2627.
 - [29] D. Amodei, S. Ananthanarayanan, R. Anubhai, J. Bai, E. Battenberg, C. Case, J. Casper, B. Catanzaro, Q. Cheng, G. Chen *et al.*, “Deep speech 2: End-to-end speech recognition in english and mandarin,” in *International conference on machine learning*, 2016, pp. 173–182.